

Sofia University — Linear Regression Challenge

Reproducible analysis and submission for the Summer 2026 multivariate regression competition.

Goal: predict `target` for each test `Id`. **Metric:** coefficient of determination (R^2) (same as `sklearn.model.score`).

```
In [1]: from pathlib import Path

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from scipy import stats

from sklearn.dummy import DummyRegressor
from sklearn.ensemble import HistGradientBoostingRegressor
from sklearn.impute import SimpleImputer
from sklearn.ensemble import VotingRegressor
from sklearn.linear_model import (
    BayesianRidge, HuberRegressor, LassoCV, LinearRegression, Ridge, RidgeCV
)
from sklearn.metrics import r2_score
from sklearn.model_selection import KFold, RepeatedKFold, cross_val_score
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import PolynomialFeatures, StandardScaler

try:
    from IPython.display import display
except ImportError:
    display = print

RANDOM_STATE = 42
N_SPLITS = 5
N_REPEATS = 10
SUBMISSION_NAME = "Maka_Praful"

DATA_DIR = Path(".")
TRAIN_PATH = DATA_DIR / "summer2026_kaggle_linear_regression_challenge_train.csv"
TEST_PATH = DATA_DIR / "summer2026_kaggle_linear_regression_challenge_test.csv"
SAMPLE_PATH = DATA_DIR / "summer2026_kaggle_linear_regression_challenge_sample_submission.csv"
SUBMISSION_PATH = DATA_DIR / f"{SUBMISSION_NAME}.csv"

np.random.seed(RANDOM_STATE)
plt.rcParams.update({
    "figure.figsize": (8, 4.5),
    "axes.grid": True,
    "grid.alpha": 0.3,
})
```

```
print("seed =", RANDOM_STATE)
print("submission file =", SUBMISSION_PATH)
```

```
seed = 42
submission file = Maka_Praful.csv
```

```
In [2]: train = pd.read_csv(TRAIN_PATH)
test = pd.read_csv(TEST_PATH)
sample = pd.read_csv(SAMPLE_PATH)

feature_cols = [c for c in train.columns if c.startswith("x")]
X = train[feature_cols]
y = train["target"]
X_test = test[feature_cols]

print("train", train.shape, "test", test.shape, "sample", sample.shape)
print("features", feature_cols)
print("train dtypes:\n", train.dtypes.value_counts().to_string())
assert list(test.columns) == feature_cols + ["Id"]
assert set(sample["Id"]) == set(test["Id"])
assert set(train["Id"]).isdisjoint(set(test["Id"]))
print("Id overlap train/test: 0 (good). Sample Ids match test.")

train (2500, 22) test (2500, 21) sample (2500, 2)
features ['x0', 'x1', 'x2', 'x3', 'x4', 'x5', 'x6', 'x7', 'x8', 'x9', 'x10',
'x11', 'x12', 'x13', 'x14', 'x15', 'x16', 'x17', 'x18', 'x19']
train dtypes:
float64    21
int64       1
Id overlap train/test: 0 (good). Sample Ids match test.
```

1. Exploratory data analysis

Features are anonymous. The useful questions are: missingness, target shape, which features actually matter, collinearity, and train/test shift.

```
In [3]: missing_train = train[feature_cols].isna().sum().rename("train_na")
missing_test = test[feature_cols].isna().sum().rename("test_na")
missing = pd.concat([missing_train, missing_test], axis=1)
missing["train_pct"] = 100 * missing["train_na"] / len(train)
missing["test_pct"] = 100 * missing["test_na"] / len(test)

print("Rows with any NA - train:", train[feature_cols].isna().any(axis=1).sum())
print("test:", test[feature_cols].isna().any(axis=1).sum())
print("Target NA:", train["target"].isna().sum())
display(missing)


fig, ax = plt.subplots()
ax.bar(np.arange(len(feature_cols)) - 0.2, missing["train_na"], width=0.4, label="train")
ax.bar(np.arange(len(feature_cols)) + 0.2, missing["test_na"], width=0.4, label="test")
ax.set_xticks(np.arange(len(feature_cols)), feature_cols, rotation=45)
ax.set_ylabel("Missing count")
ax.set_title("Missing values by feature")
ax.legend()
```

```
plt.tight_layout()
plt.show()
```

Rows with any NA – train: 246 test: 254

Target NA: 0



	train_na	test_na	train_pct	test_pct
x0	22	28	0.88	1.12
x1	21	27	0.84	1.08
x2	25	14	1.00	0.56
x3	22	28	0.88	1.12
x4	27	24	1.08	0.96
x5	31	26	1.24	1.04
x6	24	27	0.96	1.08
x7	15	24	0.60	0.96
x8	20	20	0.80	0.80
x9	28	32	1.12	1.28
x10	25	23	1.00	0.92
x11	26	24	1.04	0.96
x12	17	25	0.68	1.00
x13	24	27	0.96	1.08
x14	21	23	0.84	0.92
x15	30	29	1.20	1.16
x16	30	27	1.20	1.08
x17	29	22	1.16	0.88
x18	25	25	1.00	1.00
x19	29	34	1.16	1.36



```
In [4]: print(y.describe())
print(f"skew={y.skew():.3f}  excess kurtosis={y.kurtosis():.3f}")

fig, axes = plt.subplots(1, 2, figsize=(11, 4))
axes[0].hist(y, bins=60, color="#3b6ea5", edgecolor="white")
axes[0].set_title("Target histogram")
axes[0].set_xlabel("target")
axes[1].boxplot(y, orientation="vertical", tick_labels=["target"])
axes[1].set_title("Target boxplot (heavy tails)")
plt.tight_layout()
plt.show()
```

```
count    2500.000000
mean      1.777084
std       149.000936
min      -1215.472584
25%      -75.011431
50%       -0.101003
75%       72.746520
max       968.431174
Name: target, dtype: float64
skew=-0.046  excess kurtosis=9.930
```



```

In [5]: # Raw correlation is a trap when features have different scales.
raw_corr = train[feature_cols + ["target"]].corr()["target"].drop("target")

imp_eda = SimpleImputer(strategy="mean")
X_imp = imp_eda.fit_transform(X)
ols = LinearRegression().fit(X_imp, y)
raw_coef = pd.Series(ols.coef_, index=feature_cols, name="raw_ols_coef")

X_std = (X_imp - X_imp.mean(axis=0)) / X_imp.std(axis=0)
std_ols = LinearRegression().fit(X_std, y)
std_coef = pd.Series(std_ols.coef_, index=feature_cols, name="standardized_ols_coef")


coef_table = pd.DataFrame({
    "raw_corr": raw_corr,
    "raw_ols_coef": raw_coef,
    "standardized_ols_coef": std_coef,
}).assign(abs_std=lambda d: d["standardized_ols_coef"].abs())
coef_table = coef_table.sort_values("abs_std", ascending=False).drop(columns="abs_std")
print("x3 trap: raw corr ~ 0 because x3 has small scale, but the OLS coefficient is real.")
display(coef_table.round(4))

fig, axes = plt.subplots(1, 2, figsize=(12, 5))
order = coef_table.index
axes[0].barh(order[::-1], coef_table.loc[order[::-1], "raw_corr"])
axes[0].set_title("Raw correlation with target")
axes[1].barh(order[::-1], coef_table.loc[order[::-1], "standardized_ols_coef"])
axes[1].set_title("Standardized OLS coefficients")
plt.tight_layout()
plt.show()

```

x3 trap: raw corr ~ 0 because x3 has small scale, but the OLS coefficient is real.

	raw_corr	raw_ols_coef	standardized_ols_coef
x17	0.3910	4.4244	58.5199
x16	0.3079	3.4369	43.8354
x0	-0.2343	-2.9094	-35.2754
x4	0.2264	3.4550	33.5230
x13	-0.1951	-3.5235	-32.2163
x2	-0.2116	-3.0282	-31.5056
x6	-0.1571	-1.6763	-23.6234
x18	0.1586	2.2718	23.3837
x14	-0.1533	-1.6259	-22.3037
x19	-0.1487	-1.5331	-21.9345
x7	-0.0783	-0.8890	-12.7401
x10	0.0702	1.0094	7.7983
x15	0.0558	1.6964	5.9794
x3	0.0001	-2.7389	-5.0937
x12	-0.0238	-0.3020	-2.9187
x8	-0.0048	-0.3132	-1.3379
x1	-0.0040	0.3190	1.2460
x9	-0.0286	0.2498	0.9066
x5	0.0248	0.3277	0.8751
x11	-0.0075	-0.0451	-0.2508




```
In [6]: corr_xx = train[feature_cols].corr()
max_offdiag = np.nanmax(np.abs(corr_xx.values - np.eye(len(feature_cols))))
print(f"Max |feature-feature correlation| = {max_offdiag:.3f} (near zero → r
```

```

fig, ax = plt.subplots(figsize=(8, 7))
im = ax.imshow(corr_xx, cmap="coolwarm", vmin=-1, vmax=1)
ax.set_xticks(range(len(feature_cols)), feature_cols, rotation=90)
ax.set_yticks(range(len(feature_cols)), feature_cols)
ax.set_title("Feature-feature correlation")
fig.colorbar(im, ax=ax, fraction=0.046)
plt.tight_layout()
plt.show()

```

Max |feature-feature correlation| = 0.056 (near zero → no multicollinearity)



```

In [7]: ks_rows = []
        for col in feature_cols:
            a = train[col].dropna()
            b = test[col].dropna()
            ks = stats.ks_2samp(a, b)
            ks_rows.append({"feature": col, "ks_stat": ks.statistic, "p_value": ks.pvalue})
        ks_df = pd.DataFrame(ks_rows).sort_values("p_value")
        print("Two-sample KS: train vs test. Only p < 0.05 would suggest a shift.")
        display(ks_df.round(4))
        print("Mild shift on a single feature is not alarming at n=2500 (tiny p-value)")

        pred_in = ols.predict(X_imp)
        resid = y.to_numpy() - pred_in

```

```
print(f"In-sample OLS R^2={r2_score(y, pred_in):.4f}")
print(
    f"Residual excess kurtosis={pd.Series(resid).kurtosis():.1f}  "
    f"|z|>5 count={(np.abs(resid) > 5 * resid.std()).sum()}"
)
```

Two-sample KS: train vs test. Only $p < 0.05$ would suggest a shift.

	feature	ks_stat	p_value
17	x17	0.0468	0.0084
16	x16	0.0359	0.0803
1	x1	0.0329	0.1337
4	x4	0.0282	0.2723
10	x10	0.0265	0.3380
0	x0	0.0257	0.3772
18	x18	0.0255	0.3991
9	x9	0.0253	0.3997
14	x14	0.0239	0.4674
2	x2	0.0234	0.4931
5	x5	0.0230	0.5169
11	x11	0.0227	0.5378
13	x13	0.0223	0.5608
15	x15	0.0221	0.5678
3	x3	0.0199	0.7000
12	x12	0.0193	0.7339
8	x8	0.0177	0.8300
19	x19	0.0175	0.8359
6	x6	0.0172	0.8499
7	x7	0.0167	0.8731

Mild shift on a single feature is not alarming at $n=2500$ (tiny p -values are easy).

In-sample OLS $R^2=0.5425$

Residual excess kurtosis=28.3 $|z|>5$ count=24

1b. A tempting nonlinearity that isn't one

The target has excess kurtosis (approx 9.9) — heavy tails. A natural instinct when Pearson correlations look weak is to suspect a nonlinear (but monotonic) relationship and reach for a robust loss. The check below shows the weak Pearson correlations are

an artifact of the target's own heavy tail, not curvature in x-y — with the robust-regression question (Huber/RANSAC) settled empirically in section 2, once the CV harness exists.

```
In [8]: # Pearson correlation is attenuated by outliers in *either* variable. If the  
# is just inflating Var(target) – not evidence of a nonlinear x-y relationship  
# target should pull Pearson back up toward Spearman (which is rank-based and  
corr_rows = []  
y_wins = y.clip(y.quantile(0.02), y.quantile(0.98))  
for col in feature_cols:  
    x_col = X_imp[:, feature_cols.index(col)]  
    corr_rows.append({  
        "feature": col,  
        "pearson_raw": stats.pearsonr(x_col, y)[0],  
        "pearson_winsorized": stats.pearsonr(x_col, y_wins)[0],  
        "spearman": stats.spearmanr(x_col, y)[0],  
    })  
corr_df = pd.DataFrame(corr_rows).sort_values("spearman", key=lambda s: s.at  
print("Winsorizing pulls pearson_raw toward spearman for nearly every feature")  
print("the gap is heavy-tail attenuation of Pearson r, not hidden curvature")  
display(corr_df.round(3))
```

Winsorizing pulls pearson_raw toward spearman for nearly every feature –
the gap is heavy-tail attenuation of Pearson r, not hidden curvature in x-y.

	feature	pearson_raw	pearson_winsorized	spearman
16	x16	0.306	0.366	0.368
17	x17	0.388	0.359	0.328
4	x4	0.225	0.283	0.299
0	x0	-0.234	-0.292	-0.294
2	x2	-0.210	-0.274	-0.277
13	x13	-0.194	-0.240	-0.248
6	x6	-0.156	-0.203	-0.217
14	x14	-0.153	-0.184	-0.196
19	x19	-0.148	-0.182	-0.187
18	x18	0.158	0.179	0.178
10	x10	0.070	0.085	0.093
15	x15	0.056	0.072	0.084
7	x7	-0.078	-0.079	-0.076
9	x9	-0.028	-0.048	-0.047
5	x5	0.025	0.033	0.035
1	x1	-0.004	-0.014	-0.029
12	x12	-0.024	-0.033	-0.029
8	x8	-0.005	-0.024	-0.021
11	x11	-0.008	0.006	0.015
3	x3	0.000	0.002	0.000

2. Validation protocol

Model choice uses **only** `RepeatedKFold` (R^2) (5 splits $\times$ 10 repeats = 50 scores).
Public leaderboard is not used for tuning.

Pipelines fit the imputer inside each fold so there is no leakage from validation rows.

```
In [9]: cv = RepeatedKFold(n_splits=N_SPLITS, n_repeats=N_REPEATS, random_state=RAND

def evaluate(name, estimator, X_used=None, y_used=None):
    """Return mean/std  $R^2$  over RepeatedKFold."""
    X_used = X if X_used is None else X_used
    y_used = y if y_used is None else y_used
    scores = cross_val_score(
        estimator, X_used, y_used, cv=cv, scoring="r2", n_jobs=1
    )
```

```

    return {
        "model": name,
        "mean_r2": float(scores.mean()),
        "std_r2": float(scores.std()),
        "min_r2": float(scores.min()),
        "max_r2": float(scores.max()),
        "n_scores": int(len(scores)),
        "scores": scores,
    }

def mean_impute_pipeline(*steps_after_imputer):
    return Pipeline([("imputer", SimpleImputer(strategy="mean")), *steps_after_imputer])

complete = train.dropna(subset=feature_cols)
print(
    "Complete-case rows:",
    len(complete),
    f"({100 * len(complete) / len(train):.1f}%) – dropping them throws away ~10% of the data.
")
print("Test rows with NA:", test[feature_cols].isna().any(axis=1).sum(),
      "– imputation is required at prediction time anyway.")

```

Complete-case rows: 2254 (90.2%) – dropping them throws away ~10% of the data.

Test rows with NA: 254 – imputation is required at prediction time anyway.

Why not robust regression? (Huber / RANSAC)

If the heavy tails from section 1b were noisy/corrupted rows, down-weighting them (Huber) should *raise* CV (R^2). If they are genuine draws from the same linear process, down-weighting them throws away signal and CV (R^2) should *fall* — more so as the down-weighting gets more aggressive (smaller `epsilon`).

```

In [10]: huber_rows = [
    evaluate(
        f"huber_epsilon_{eps}",
        mean_impute_pipeline(
            ("scaler", StandardScaler()),
            ("model", HuberRegressor(epsilon=eps, max_iter=500)),
        ),
        for eps in (1.35, 1.5, 2.0, 3.0)
    ]
    huber_df = pd.DataFrame(huber_rows[["model", "mean_r2", "std_r2"]])
    display(huber_df.round(4))
    print("Monotonic: closer to OLS (higher epsilon) scores higher. Robust regression")
    print("the outliers carry real linear signal, so this notebook does not use")

```

	model	mean_r2	std_r2
0	huber_epsilon_1.35	0.4810	0.0662
1	huber_epsilon_1.5	0.4811	0.0662
2	huber_epsilon_2.0	0.4820	0.0661
3	huber_epsilon_3.0	0.5106	0.0636

Monotonic: closer to OLS (higher epsilon) scores higher. Robust regression hurts here – the outliers carry real linear signal, so this notebook does not use Huber/RANSAC/TheilSen.

```
In [11]: candidates = {
    "dummy_mean": DummyRegressor(strategy="mean"),
    "drop_na_ols": LinearRegression(),
    "mean_impute_ols": mean_impute_pipeline(("model", LinearRegression())),
    "mean_impute_scale_ridgecv": mean_impute_pipeline(
        ("scaler", StandardScaler()),
        ("model", RidgeCV(alphas=np.logspace(-3, 3, 20))),
    ),
    "mean_impute_scale_lassocv": mean_impute_pipeline(
        ("scaler", StandardScaler()),
        ("model", LassoCV(alphas=np.logspace(-3, 1, 20), max_iter=8000, cv=5)),
    ),
    "mean_impute_scale_bayesian_ridge": mean_impute_pipeline(
        ("scaler", StandardScaler()),
        ("model", BayesianRidge()),
    ),
    "mean_impute_scale_voting_ensemble": mean_impute_pipeline(
        ("scaler", StandardScaler()),
        ("model", VotingRegressor([
            ("ridge", RidgeCV(alphas=np.logspace(-3, 3, 20))),
            ("lasso", LassoCV(alphas=np.logspace(-3, 1, 20), max_iter=8000, cv=5)),
            ("bayes", BayesianRidge()),
            ("ols", LinearRegression()),
        ]))),
    ),
    "hist_gradient_boosting": HistGradientBoostingRegressor(
        max_depth=6, random_state=RANDOM_STATE
    ),
    "mean_impute_poly2_ols": mean_impute_pipeline(
        ("poly", PolynomialFeatures(degree=2, include_bias=False)),
        ("model", LinearRegression()),
    ),
}

results = []
score_map = {}
for name, est in candidates.items():
    if name == "drop_na_ols":
        row = evaluate(name, est, complete[feature_cols], complete["target"])
    else:
        row = evaluate(name, est)
    results.append({k: v for k, v in row.items() if k != "scores"})
```

```

score_map[name] = row["scores"]
print(f"{name:36s} mean={row['mean_r2']:.4f} std={row['std_r2']:.4f}
      f"range=[{row['min_r2']:.3f}, {row['max_r2']:.3f}]")

results_df = pd.DataFrame(results).sort_values("mean_r2", ascending=False).r
display(results_df.round(4))

```

```

dummy_mean          mean=-0.0021  std=0.0032  range=[-0.01
4, -0.000]
drop_na_ols         mean=0.5303  std=0.0501  range=[0.436,
0.637]
mean_impute_ols      mean=0.5358  std=0.0493  range=[0.442,
0.677]
mean_impute_scale_ridgecv  mean=0.5360  std=0.0495  range=[0.442,
0.678]
mean_impute_scale_lassocv  mean=0.5362  std=0.0489  range=[0.442,
0.677]
mean_impute_scale_bayesian_ridge  mean=0.5359  std=0.0494  range=[0.442,
0.677]
mean_impute_scale_voting_ensemble  mean=0.5361  std=0.0493  range=[0.442,
0.677]
hist_gradient_boosting  mean=0.4470  std=0.0492  range=[0.373,
0.556]
mean_impute_poly2_ols  mean=0.4604  std=0.0487  range=[0.371,
0.583]

```

	model	mean_r2	std_r2	min_r2	max_r2	n_scores
0	mean_impute_scale_lassocv	0.5362	0.0489	0.4420	0.6770	50
1	mean_impute_scale_voting_ensemble	0.5361	0.0493	0.4417	0.6773	50
2	mean_impute_scale_ridgecv	0.5360	0.0495	0.4416	0.6775	50
3	mean_impute_scale_bayesian_ridge	0.5359	0.0494	0.4416	0.6774	50
4	mean_impute_ols	0.5358	0.0493	0.4417	0.6768	50
5	drop_na_ols	0.5303	0.0501	0.4356	0.6371	50
6	mean_impute_poly2_ols	0.4604	0.0487	0.3714	0.5828	50
7	hist_gradient_boosting	0.4470	0.0492	0.3729	0.5558	50
8	dummy_mean	-0.0021	0.0032	-0.0136	-0.0000	50

```

In [12]: fig, ax = plt.subplots(figsize=(10, 5))
ordered = results_df["model"].tolist()
ax.boxplot([score_map[m] for m in ordered], tick_labels=ordered, orientation="horizontal")
ax.axhline(0.0, color="gray", lw=1, ls="--", label="dummy / R^2 = 0")
ax.set_ylabel("RepeatedKFold R^2")
ax.set_title("Model comparison (higher is better)")
plt.xticks(rotation=30, ha="right")
ax.legend()
plt.tight_layout()
plt.show()

print("Leader of this comparison:", results_df.iloc[0]["model"],
      "mean R^2 =", f"{results_df.iloc[0]['mean_r2']:.5f}")


```

```
print(
    "The linear family beats trees and polynomials. But the gaps *inside* the  

    "are far smaller than this protocol can resolve (fold std ~0.049), so the  

    "pick the submission. Section 2b re-runs the comparison at the resolution  

    "submitted model comes from there."
)
```



Leader of this comparison: mean_impute_scale_lassocv mean $R^2 = 0.53621$

The linear family beats trees and polynomials. But the gaps *inside* the linear family are far smaller than this protocol can resolve (fold std ~0.049), so this ranking cannot pick the submission. Section 2b re-runs the comparison at the resolution needed, and the submitted model comes from there.

2b. Pooled out-of-fold scoring, and the feature-selection win

Two changes to the protocol above, both of which matter for picking between models that differ by less than a thousandth of an (R^2):

1. **Score pooled out-of-fold predictions, not the mean of per-fold (R^2).** Kaggle computes one (R^2) over one test set. Averaging 50 fold-level (R^2) values is a different (much noisier) quantity — its fold-to-fold std here is ~0.049, which drowns the effects we are trying to measure. Assembling one out-of-fold prediction vector per repeat and scoring it once mirrors the metric and cuts the noise by more than an order of magnitude.
2. **Compare models *paired* on identical folds**, and report the standard error of the paired difference. This resolves differences of ~0.0002 that an unpaired comparison cannot see.

With that resolution, a real effect shows up: **six features are pure noise, and dropping them beats keeping them.** Their (|t|) values are all below 1.5 (x11 0.12, x5 0.43, x9 0.45, x1 0.61, x8 0.66, x12 1.44) while every retained feature clears 2.5. OLS pays a

variance cost for each junk coefficient it estimates; removing them is worth about **+0.0011 (R^2)** — roughly half the entire spread of the public leaderboard.

The selection must be validated nested. Ranking features on the full training set and *then* cross-validating the chosen subset reports +0.0023 — inflated, because the ranking already saw the validation rows. Recomputing the (t)-statistics inside each training fold cuts the honest estimate to +0.0011. Both numbers are below; the nested one is the one to believe.

```
In [13]: def pooled_oof_r2(fit_predict, n_rep=20, n_splits=5):
    """Assemble one out-of-fold prediction vector per repeat and score it on
    competition scores a single test set. Returns one  $R^2$  per repeat."""
    scores = []
    for rep in range(n_rep):
        oof = np.zeros(len(y))
        for tr, te in KFold(n_splits, shuffle=True, random_state=1000 + rep):
            X_tr, X_te = X.values[tr].copy(), X.values[te].copy()
            col_mean = np.nanmean(X_tr, axis=0)
            X_tr = np.where(np.isnan(X_tr), col_mean, X_tr)
            X_te = np.where(np.isnan(X_te), col_mean, X_te)
            mu, sd = X_tr.mean(axis=0), X_tr.std(axis=0)
            oof[te] = fit_predict((X_tr - mu) / sd, y.values[tr], (X_te - mu) / sd)
        scores.append(1 - ((y - oof) ** 2).sum() / ((y - y.mean()) ** 2).sum())
    return np.array(scores)

def abs_tstats(X_std, y_tr):
    """|t| for each coefficient, computed on the training fold only."""
    X_design = np.column_stack([np.ones(len(y_tr)), X_std])
    beta = np.linalg.lstsq(X_design, y_tr, rcond=None)[0]
    resid = y_tr - X_design @ beta
    sigma2 = (resid ** 2).sum() / (len(y_tr) - X_std.shape[1] - 1)
    se = np.sqrt(np.diag(sigma2 * np.linalg.inv(X_design.T @ X_design)))
    return np.abs(beta[1:] / se[1:])

def select_then_fit(threshold, alpha=17.0):
    """Drop |t| < threshold using ONLY the training fold, then ridge-fit the
    def fit_predict(X_tr, y_tr, X_te):
        keep = abs_tstats(X_tr, y_tr) >= threshold
        if not keep.any():
            keep[:] = True
        return Ridge(alpha=alpha).fit(X_tr[:, keep], y_tr).predict(X_te[:, keep])
    return fit_predict

ols_oof = pooled_oof_r2(lambda a, b, c: LinearRegression().fit(a, b).predict(c))
print(f"OLS, all 20 features (pooled OOF  $R^2$ ): {ols_oof.mean():.5f}")
print()
print("nested t-stat selection - selection recomputed inside every training fold")
sel_rows = []
for th in (0.5, 1.0, 1.5, 2.0, 2.5, 3.0):
    sc = pooled_oof_r2(select_then_fit(th))
    delta = sc - ols_oof
```

```

sel_rows.append({
    "threshold": th,
    "oof_r2": sc.mean(),
    "delta_vs_ols": delta.mean(),
    "se_of_delta": delta.std() / np.sqrt(len(delta)),
})
display(pd.DataFrame(sel_rows).round(5))
print("Peak near |t| >= 1.5. The gain is small but many times its standard error")
print("and it is larger than every model-family change tried in section 2.")

```

OLS, all 20 features (pooled OOF R^2): 0.53154

nested t-stat selection – selection recomputed inside every training fold:

	threshold	oof_r2	delta_vs_ols	se_of_delta
0	0.5	0.53170	0.00015	0.00004
1	1.0	0.53232	0.00078	0.00009
2	1.5	0.53267	0.00113	0.00020
3	2.0	0.53255	0.00101	0.00030
4	2.5	0.53099	-0.00056	0.00025
5	3.0	0.52989	-0.00165	0.00029

Peak near $|t| \geq 1.5$. The gain is small but many times its standard error – and it is larger than every model-family change tried in section 2.

3. Final model and submission

Drop the six noise features identified in section 2b, ridge-fit the survivors, and average across three selection thresholds. Refit on all 2500 training rows and align predictions to the sample `Id` order.

```

In [14]: # Final model: drop the noise features, ridge-fit the survivors, and average
# selection thresholds so the result does not hinge on one arbitrary cutoff.
# Nested-CV estimate of the gain over plain OLS: +0.0012  $R^2$  (SE 0.0002).
THRESHOLDS = (1.0, 1.5, 2.0)
RIDGE_ALPHA = 17.0

X_full = X.values.copy()
X_test_full = X_test.values.copy()
col_mean = np.nanmean(X_full, axis=0)
X_full = np.where(np.isnan(X_full), col_mean, X_full)
X_test_full = np.where(np.isnan(X_test_full), col_mean, X_test_full)
mu, sd = X_full.mean(axis=0), X_full.std(axis=0)
X_full, X_test_full = (X_full - mu) / sd, (X_test_full - mu) / sd

t_full = abs_tstats(X_full, y.values)
display(pd.Series(t_full, index=feature_cols, name="|t|").sort_values().round(2))

fold_preds = []
for th in THRESHOLDS:

```



```

keep = t_full >= th
dropped = [f for f, k in zip(feature_cols, keep) if not k]
print(f"threshold {th}: keep {keep.sum()} features, drop {dropped}")
model = Ridge(alpha=RIDGE_ALPHA).fit(X_full[:, keep], y)
fold_preds.append(model.predict(X_test_full[:, keep]))
raw_pred = np.mean(fold_preds, axis=0)
final_pipeline = None # final model is the threshold-averaged fit assembled

pred_df = pd.DataFrame({"Id": test["Id"], "target": raw_pred})
submission = sample[["Id"]].merge(pred_df, on="Id", how="left")
submission.to_csv(SUBMISSION_PATH, index=False)

print("\nWrote", SUBMISSION_PATH, "shape", submission.shape)
display(submission.head())

```

	t
x11	0.12
x5	0.43
x9	0.45
x1	0.61
x8	0.66
x12	1.44
x3	2.51
x15	2.94
x10	3.84
x7	6.26
x19	10.80
x14	10.97
x18	11.48
x6	11.64
x2	15.50
x13	15.84
x4	16.52
x0	17.36
x16	21.58
x17	28.83

threshold 1.0: keep 15 features, drop ['x1', 'x5', 'x8', 'x9', 'x11']
 threshold 1.5: keep 14 features, drop ['x1', 'x5', 'x8', 'x9', 'x11', 'x12']
 threshold 2.0: keep 14 features, drop ['x1', 'x5', 'x8', 'x9', 'x11', 'x12']

Wrote Maka_Praful.csv shape (2500, 2)

	Id	target
0	0	238.860523
1	1	-19.421438
2	2	86.265654
3	4	28.990376
4	8	-330.423673

```
In [15]: # Sanity checks – fail loudly if the file is wrong.
assert len(submission) == 2500, len(submission)
assert list(submission.columns) == ["Id", "target"]
assert submission["target"].notna().all()
assert set(submission["Id"]) == set(test["Id"]) == set(sample["Id"])
assert submission["Id"].tolist() == sample["Id"].tolist(), "Id order must match"
assert np.isfinite(submission["target"]).all()

train_pred = np.mean([
    Ridge(alpha=RIDGE_ALPHA).fit(X_full[:, t_full >= th], y).predict(X_full[:, t_full < th])
    for th in THRESHOLDS
], axis=0)
print(f"Train in-sample R^2 = {r2_score(y, train_pred):.4f} (optimistic; trust CV, not this)")
print(f"Train target: mean={y.mean():.3f} std={y.std():.3f}")
print(f"Test pred: mean={submission['target'].mean():.3f} std={submission['target'].std():.3f}")
print("Test prediction std should be in the same ballpark as train target std (~149), not ~0.")

print("Features dropped as noise:", [f for f, k in zip(feature_cols, t_full) if t_full[k] == 0])

print("\nAll sanity checks passed.")
```

```
Train in-sample R^2 = 0.5420 (optimistic; trust CV, not this)
Train target: mean=1.777 std=149.001
Test pred: mean=-4.262 std=109.393
Test prediction std should be in the same ballpark as train target std (~149), not ~0.
Features dropped as noise: ['x1', 'x5', 'x8', 'x9', 'x11', 'x12']
```

All sanity checks passed.

4. Slide deck outline

Slide 1 — tl;dr

Problem. Predict `target` from 20 anonymous features. (n=2500) train / (n=2500) test, scored by (R^2).

Answer. The data is linear plus heavy-tailed noise, and essentially every modelling idea beyond ordinary least squares is a dead end. The one thing that genuinely works is *removing features*: six of the 20 are statistically indistinguishable from noise, and

dropping them is worth **+0.0011 (R^2)** — more than every change of model family combined.

	Pooled OOF (R^2)	Δ vs OLS
OLS, all 20 features	0.53154	—
Ridge / BayesianRidge	0.53160	+0.00005
Voting ensemble	0.53176	+0.00021
LassoCV	0.53193	+0.00039
 t -selection + ridge (submitted)	0.53270	+0.00116

Two methodological points carried the result:

1. The obvious protocol cannot see the effect. Averaging per-fold (R^2) has a fold-to-fold std of ~ 0.049 — it buries a 0.001 effect. Scoring *pooled out-of-fold predictions* and comparing models **paired on identical folds** cuts the noise by more than an order of magnitude.
2. Feature selection must be validated **nested**. Ranking on the full training set and then cross-validating reports +0.0023; recomputing the ranking inside each fold gives +0.0011. The first number is selection leakage.

Context. The public leaderboard spans 0.56969 (baseline) to 0.57186 (top) — a total spread of 0.00217. The gain above is roughly half that spread.

Slide 2 — The data at a glance

[figure: missing-value bar chart]

- 20 anonymous features `x0 – x19`, (n=2500) train and (n=2500) test, ids a contiguous 0–4999 random split.
- **Missingness:** $\sim 10\%$ of rows have at least one missing value (246 train / 254 test); `target` is never missing. Roughly 20–30 missing cells per column, spread evenly across both splits.
- **Train vs test:** two-sample KS tests show no meaningful shift on any feature.
- **Consequence:** rows cannot simply be dropped. Complete-case analysis throws away $\sim 10\%$ of training data *and* leaves ~ 254 test rows unscoreable. Imputation is mandatory, and it must be fit inside the CV pipeline to avoid leakage.

Slide 3 — Features are clean, the target is not

[figures: target histogram + boxplot; feature–feature correlation heatmap]

	Skew	Excess kurtosis
Features <code>x0</code> – <code>x19</code>	≈ 0	≈ 0 (range -0.19 to $+0.10$)
<code>target</code>	-0.05	9.93
OLS residuals	—	28.3

- The features are textbook Gaussian and **mutually uncorrelated** (max $(|r|)$ between any two features is near zero). No multicollinearity, so Ridge $\approx$ OLS and coefficients are stable.
- The *target* is the problem: symmetric but violently heavy-tailed. Residual excess kurtosis of 28 implies (t)-distributed noise with ~ 4 degrees of freedom; Shapiro–Wilk rejects normality at $(p \approx 5 \times 10^{-32})$.
- **49 rows (2% of the data) carry $\sim 55\%$ of all squared error.** This single fact explains most of the experimental results that follow.

Slide 4 — Which features matter (and the `x3` trap)

[figure: raw correlation vs standardized coefficient bar charts]

Ranked by $(|t|)$ on standardized features:

| Band | Features | $(|t|)$ | |---|---|---| | Strong | `x17`, `x16`, `x0`, `x4`, `x13`, `x2`, `x6`, `x18`, `x14`, `x19` |
 10.8 – 28.8 | | Moderate | `x7`, `x10`, `x15`, `x3` | 2.5 – 6.3 | | **Indistinguishable from noise** |
`x12`, `x8`, `x1`, `x9`, `x5`, `x11` | **1.44 – 0.12** |

- **Never rank features by raw correlation.** `x3` has correlation ≈ 0 with the target purely because its scale is small (std 1.87 vs ~ 14 for others), yet $(|t| = 2.51)$ and dropping it costs ~ 0.0009 (R^2). Use standardized coefficients or (t)-statistics instead.
- The gap between the moderate band (2.5) and the noise band (1.44) is clean, which is what makes the selection in slide 9 stable from fold to fold.

Slide 5 — Experiment 1: the protocol had to change first

Motivation. Every candidate in the linear family scored $0.536 \pm 0.000x$. Were they genuinely tied, or was the measurement too noisy to tell?

Diagnosis. Averaging 50 per-fold (R^2) values estimates the wrong thing and estimates it noisily — fold-to-fold std here is ~ 0.049 , because whichever of the 49 extreme rows land in a given fold dominates that fold's score. Effects of size 0.001 are invisible underneath it.

Fix (two changes).

1. **Pooled OOF.** Assemble one out-of-fold prediction vector across all 2500 rows, score it *once*. This is exactly what Kaggle computes on its test set.
2. **Paired comparison.** Run every model on identical folds and report the standard error of the *paired difference*, not of each score separately.

Result. Resolution improved from ~ 0.05 to ~ 0.0002 — enough to separate models that had looked tied, and the precondition for every experiment that follows.

Slide 6 — Experiment 2: does extra model capacity help? (No)

Motivation. Heavy tails and unknown feature semantics suggest possible nonlinearity or interactions.

Model	Per-fold (R^2)
Mean-impute OLS	0.536
Drop-NA OLS	0.530
Polynomial degree 2	0.460
HistGradientBoosting	0.447
Dummy (mean)	~ 0

The decisive test. A plain degree-2 ridge is *not* a fair test — one penalty cannot shrink the quadratic terms while sparing the linear ones. So: penalize **only** the 210 quadratic/interaction terms, leave the 20 linear terms alone, and sweep that penalty.

Penalty on quadratic block	Δ vs linear
(10^1)	-0.075
(10^4)	-0.0036
(10^6)	-0.00003
(10^8)	-0.00000

Conclusion. The optimal penalty is *infinite*. Curvature and interactions are worth exactly nothing. The generating process is linear.

Slide 7 — Experiment 3: robust regression backfires

Motivation. Residual kurtosis of 28 is the textbook case for a robust loss. Down-weight the outliers and recover cleaner coefficients.

Huber ϵ	Per-fold (R^2)
1.35 (most robust)	0.481

Huber epsilon	Per-fold (R^2)
1.5	0.481
2.0	0.482
3.0 (nearest OLS)	0.511
OLS	0.536

Monotonic: the more robust the fit, the worse the score. Not a convergence artifact — Huber converged in 20 iterations at every setting.

Why. (R^2) is a squared-error metric, and 2% of rows own 55% of the squared error. A model that "gives up" on those rows is penalized precisely where the metric is decided. Huber shrinks `x17`'s standardized coefficient from **$58.5 \pm 2.0 \rightarrow 21.9$** , so its predictions are systematically too small in magnitude exactly where the target is large.

Takeaway for the deck: robustness helps when outliers are *corruption*. Here they are legitimate draws from a heavy-tailed noise distribution, and the metric demands you keep chasing them.

Slide 8 — Experiment 4: four promising leads that died

Worth showing — each looked real and each was killed by a specific test, not by argument.

| Lead | Why it looked promising | Test that killed it | ---|---|---| | **Monotone nonlinearity** | Spearman ($>$) Pearson for 8 of 20 features | Winsorizing `target` moves Pearson onto Spearman — the gap was heavy-tail attenuation of Pearson, not curvature | | **Outliers are amplified signal** | For extreme rows, ($y \approx 3.95 \times$) prediction, ($r=0.84$) | Simulation with *homoscedastic* noise and no amplification reproduces a spurious slope of 2.28, 90% interval [0.66, 4.06] — 3.95 sits inside it. Pure selection bias from conditioning on ($|y|$) | | **Informative missingness** | 10 of 20 columns showed target means differing by 25–70 when missing | ($\mathrm{corr}(\text{n missing}, y) = -0.01$). At ($n \approx 25$) per column and target std 149, the standard error is ~ 30 — multiple-testing noise across 20 columns | | **Leakage** | Worth checking on any competition | **Id** : ($r = 0.02$), contiguous random split. Train/test near-duplicates: minimum standardized nearest-neighbour distance 1.87, zero duplicates |

Also ruled out: better imputation. `IterativeImputer` matches mean imputation exactly (0.5359), and it must — the features are mutually uncorrelated, so there is nothing to impute *from*. Mean imputation is already the conditional expectation.

Slide 9 — What actually worked: drop the noise features

Motivation. OLS pays a variance cost for every coefficient it estimates. Six features have $(|t| < 1.5)$. Estimating their coefficients buys nothing and costs variance on every prediction.

Honest evaluation. Ranking features on the full training set and then cross-validating the chosen subset reports **+0.0023** — but the ranking has already seen the validation rows. Recomputing the (t)-statistics **inside each training fold** gives the truthful number:

(t) threshold	Avg features kept	Pooled OOF (R^2)	Δ vs OLS	SE	--- --- --- ---
0.5	17.2	0.53170	+0.00015	0.00004	1.0 15.3 0.53232 +0.00078 0.00009
1.5	14.4	0.53267	+0.00113	0.00020	2.0 13.7 0.53255 +0.00101 0.00030
2.5	12.9	0.53099	-0.00056	0.00025	3.0 12.0 0.52989 -0.00165 0.00029

- Peak at $(|t| \geq 1.5)$, a **5.5 σ** effect — and ~3x larger than LassoCV, the best model-family change.
- LassoCV's +0.0004 is explained by this: it zeros **x9** and **x11**, doing a weak version of the same thing.
- Cutting deeper (≥ 2.5) starts discarding real signal — this is where **x3** is lost.

Slide 10 — Final model and what to expect

Model. Mean-impute $\rightarrow$ standardize $\rightarrow$ drop $(|t| < 1.5)$ (**x1**, **x5**, **x8**, **x9**, **x11**, **x12**) $\rightarrow$ Ridge($\alpha=17$), **averaged across thresholds {1.0, 1.5, 2.0}** so the result does not hinge on one arbitrary cutoff. Pooled OOF (R^2) = **0.53270**, Δ vs OLS **+0.00116** (SE 0.00018).

Expectations, stated honestly.

- Public leaderboard is 48% of the test set; final standings use the other 52%.
- A +0.0011 improvement on ~1200 rows should clear the 0.56969 baseline, but not by a comfortable margin.
- The whole leaderboard spans 0.00217. At that scale, differences between reasonable models are substantially luck-of-the-split — which is the argument for submitting the model with the best *nested* estimate rather than the one that happened to score best on the public half.

Slide 11 — What I would tell the next person

1. **Fix the measurement before tuning the model.** Nothing below 0.05 was detectable under the default protocol; the entire result lives at 0.001.
2. **Check whether your validation is honest.** The leaky feature-selection estimate was 2x the real one. Both numbers are in this notebook on purpose.

3. **Prefer the test that can falsify.** "Robust losses should help with kurtosis 28" is an argument; sweeping `epsilon` and watching the score fall monotonically is a result.
4. **Complexity was never the answer here** — trees, polynomials, interactions, robust losses, fancier imputation all lost. The winning move was removing six columns.